

Building images



This chapter covers

- Some basics of image creation
- Manipulating the Docker build cache for fast and reliable builds
- Configuring timezones as part of an image build
- Running commands directly on your containers from the host
- Drilling down into the layers created by an image build
- Using the more advanced ONBUILD feature when building and using images

To get beyond the basics of using Docker, you'll want to start creating your own building blocks (images) to pull together in interesting ways. This chapter will cover some of the important parts of image creation, looking at practicalities that you might otherwise stumble over.

4.1 Building images

Although the simplicity of Dockerfiles makes them a powerful time-saving tool, there are some subtleties that can cause confusion. We'll take you over a few time-saving

features and their details, starting with the `ADD` instruction. Then we'll cover the Docker build cache, how it can let you down, and how you can manipulate it to your advantage.

Remember to refer to the official Docker documentation for complete Dockerfile instructions at <https://docs.docker.com>.

TECHNIQUE 20 Injecting files into your image using `ADD`

Although it's possible to add files within a Dockerfile using the `RUN` command and basic shell primitives, this can quickly become unmanageable. The `ADD` command was added to the list of Dockerfile commands to address the need to put large numbers of files into an image without fuss.

PROBLEM

You want to download and unpack a tarball into your image in a concise way.

SOLUTION

Tar and compress your files, and use the `ADD` directive in your Dockerfile.

Create a fresh environment for this Docker build with `mkdir add_example && cd add_example`. Then retrieve a tarball and give it a name you can reference later.

Listing 4.1 Downloading a TAR file

```
$ curl \
https://www.flamingspork.com/projects/libeatmydata/
➡ libeatmydata-105.tar.gz > my.tar.gz
```

In this case we've used a TAR file from another technique, but it could be any tarball you like.

Listing 4.2 Adding a TAR file to an image

```
FROM debian
RUN mkdir -p /opt/libeatmydata
ADD my.tar.gz /opt/libeatmydata/
RUN ls -lRt /opt/libeatmydata
```

Build this Dockerfile with `docker build --no-cache .` and the output should look like this:

Listing 4.3 Building an image with a TAR file

```
$ docker build --no-cache .
Sending build context to Docker daemon 422.9 kB
Sending build context to Docker daemon
Step 0 : FROM debian
----> c90d655b99b2
Step 1 : RUN mkdir -p /opt/libeatmydata
----> Running in fe04bac7df74
----> c0ab8c88bb46
Removing intermediate container fe04bac7df74
```

```

Step 2 : ADD my.tar.gz /opt/libeatmydata/
--> 06dcd7a88eb7
Removing intermediate container 3f093a1f9e33
Step 3 : RUN ls -lRt /opt/libeatmydata
--> Running in e3283848ad65
/opt/libeatmydata:
total 4
drwxr-xr-x 7 1000 1000 4096 Oct 29 23:02 libeatmydata-105

/opt/libeatmydata/libeatmydata-105:
total 880
drwxr-xr-x 2 1000 1000 4096 Oct 29 23:02 config
drwxr-xr-x 3 1000 1000 4096 Oct 29 23:02 debian
drwxr-xr-x 2 1000 1000 4096 Oct 29 23:02 docs
drwxr-xr-x 3 1000 1000 4096 Oct 29 23:02 libeatmydata
drwxr-xr-x 2 1000 1000 4096 Oct 29 23:02 m4
-rw-r--r-- 1 1000 1000 9803 Oct 29 23:01 config.h.in
[...edited...]
-rw-r--r-- 1 1000 1000 1824 Jun 18 2012 pandora_have_better_malloc.m4
-rw-r--r-- 1 1000 1000 742 Jun 18 2012 pandora_header_assert.m4
-rw-r--r-- 1 1000 1000 431 Jun 18 2012 pandora_version.m4
--> 2ee9b4c8059f
Removing intermediate container e3283848ad65
Successfully built 2ee9b4c8059f

```

You can see from this output that the tarball has been unpacked into the target directory by the Docker daemon (the extended output of all the files has been edited). Docker will unpack tarfiles of most standard types (.gz, .bz2, .xz, .tar).

It's worth observing that although you can download tarballs from URLs, they'll only be unpacked automatically if they're stored in the local filesystem. This can lead to confusion.

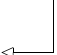
If you repeat the preceding process with the following Dockerfile, you'll notice that the file is downloaded but not unpacked.

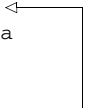
Listing 4.4 Directly adding of the TAR file from the URL

```

FROM debian
RUN mkdir -p /opt/libeatmydata
ADD \
  https://www.flamingspork.com/projects/libeatmydata/libeatmydata-105.tar.gz \
  /opt/libeatmydata/
  RUN ls -lRt /opt/libeatmydata

```


The file is retrieved from the internet using a URL.


The destination directory is indicated by the directory name and a trailing slash. Without the trailing slash, the argument is treated as a filename for the downloaded file.

Here's the resulting build output:

```

Sending build context to Docker daemon 422.9 kB
Sending build context to Docker daemon
Step 0 : FROM debian
--> c90d655b99b2

```

```

Step 1 : RUN mkdir -p /opt/libeatmydata
----> Running in 6ac454c52962
----> bdd948e413c1
Removing intermediate container 6ac454c52962
Step 2 : ADD \
https://www.flamingospork.com/projects/libeatmydata/libeatmydata-105.tar.gz
➡ /opt/libeatmydata/
Downloading [=====] \
419.4 kB/419.4 kB
----> 9d8758e90b64
Removing intermediate container 02545663f13f
Step 3 : RUN ls -lRt /opt/libeatmydata
----> Running in a947eaa04b8e
/opt/libeatmydata:
total 412
-rw----- 1 root root 419427 Jan  1  1970 \
libeatmydata-105.tar.gz
----> f18886c2418a
Removing intermediate container a947eaa04b8e
Successfully built f18886c2418a

```

The libeatmydata-105.tar.gz file has been downloaded and placed in the /opt/libeatmydata directory without being unpacked.

Note that without the trailing slash in the ADD line in the previous Dockerfile, the file would be downloaded and saved with that filename. The trailing slash indicates that the file should be downloaded and placed in the directory specified.

All new files and directories are owned by root (or whoever has group or user IDs of 0 within the container).

Whitespace in filenames

If your filenames have whitespace in them, you'll need to use the quoted form of ADD (or COPY):

```
ADD "space file.txt" "/tmp/space file.txt"
```

DISCUSSION

The ADD Dockerfile instruction is quite a workhorse, with a number of different pieces of functionality you can take advantage of. If you're going to write more than a couple of Dockerfiles (which you likely will as you go through this book), it's worth reading the official Dockerfile instructions documentation—there aren't many (18 instructions are listed in the documentation at the time of writing at <https://docs.docker.com/engine/reference/builder>) and you'll only use a few of them regularly.

People often ask about adding compressed files without extracting them. For this you should use the COPY command, which looks exactly like the ADD command but doesn't unpack any files and won't download over the internet.

TECHNIQUE 21 Rebuilding without the cache

Building with Dockerfiles takes advantage of a useful caching feature: steps that have already been built are only rebuilt if the commands have changed. The next listing shows the output of a rebuild of the to-do app from chapter 1.

Listing 4.5 Rebuilding with the cache

```
$ docker build .
Sending build context to Docker daemon 2.56 kB
Sending build context to Docker daemon
Step 0 : FROM node
---> 91cbcf796c2c
Step 1 : MAINTAINER ian.miell@gmail.com
---> Using cache
---> 8f5a8a3d9240
Step 2 : RUN git clone -q https://github.com/docker-in-practice/todo.git
---> Using cache
---> 48db97331aa2
Step 3 : WORKDIR todo
---> Using cache
---> c5c85db751d6
Step 4 : RUN npm install > /dev/null
---> Using cache
---> be943c45c55b
Step 5 : EXPOSE 8000
---> Using cache
---> 805b18d28a65
Step 6 : CMD npm start
---> Using cache
---> 19525d4ec794
Successfully built 19525d4ec794
```

Indicates you're using the cache

Specifies the cached image/layer ID

The final image is "rebuilt," but in reality nothing has changed.

As useful and time-saving as this is, it's not always the behavior you want.

Taking the preceding Dockerfile as an example, imagine you'd changed your source code and pushed it to the Git repository. The new code wouldn't be checked out, because the `git clone` command hasn't changed. As far as the Docker build is concerned, it's the same, so the cached image can be reused.

In these cases, you'll want to rebuild your image without using the cache.

PROBLEM

You want to rebuild your Dockerfile without using the cache.

SOLUTION

To force a rebuild without using the image cache, run your `docker build` with the `--no-cache` flag. The following listing runs the previous build with `--no-cache`.

Listing 4.6 Forcing a rebuild without using the cache

```
$ docker build --no-cache .
Sending build context to Docker daemon 2.56 kB
Sending build context to Docker daemon
```

Rebuilds the Docker image, ignoring cached layers with the `--no-cache` flag

```

Step 0 : FROM node
----> 91cbcf796c2c
Step 1 : MAINTAINER ian.miell@gmail.com
----> Running in ca243b77f6a1
----> 602f1294d7f1
Removing intermediate container ca243b77f6a1
Step 2 : RUN git clone -q https://github.com/docker-in-practice/todo.git
----> Running in f2c0ac021247
----> 04ee24faaf18
Removing intermediate container f2c0ac021247
Step 3 : WORKDIR todo
----> Running in c2d9cd32c182
----> 4e0029de9074
Removing intermediate container c2d9cd32c182
Step 4 : RUN npm install > /dev/null
----> Running in 79122dbf9e52
npm WARN package.json todomvc-swarm@0.0.1 No repository field.
----> 9b6531f2036a
Removing intermediate container 79122dbf9e52
Step 5 : EXPOSE 8000
----> Running in d1d58e1c4b15
----> f7c1b9151108
Removing intermediate container d1d58e1c4b15
Step 6 : CMD npm start
----> Running in 697713ebb185
----> 74f9ad384859
Removing intermediate container 697713ebb185
Successfully built 74f9ad384859

```

No mention of caching this time

Intervening images have a different ID than in the previous listing.

A new image is built.

The output shows no mention of caching, and each intervening layer ID is different from the output in listing 4.5.

Similar problems can occur in other situations. We were flummoxed early on using Dockerfiles when a network blip meant that a command didn't retrieve something properly from the network, but the command didn't error. We kept calling `docker build`, but the resulting bug wouldn't go away! This was because a "bad" image had found its way into the cache, and we didn't understand the way Docker caching worked. Eventually we figured it out.

DISCUSSION

Removing caching can be a useful sanity check once you've got your final Dockerfile, to make sure it works from top to bottom, particularly when you're using internal web resources in your company that you may have changed while iterating on the Dockerfile. This situation doesn't occur if you're using `ADD`, because Docker will download the file every time to check if it has changed, but that behavior can be tiresome if you're pretty sure it's going to stay the same and you just want to get going with writing the rest of the Dockerfile.

TECHNIQUE 22 Busting the cache

Using the `--no-cache` flag is often enough to get around any problems with the cache, but sometimes you'll want a more fine-grained solution. If you have a build that

takes a long time, for example, you may want to use the cache up to a certain point, and then invalidate it to rerun a command and create a new image.

PROBLEM

You want to invalidate the Docker build cache from a specific point in the Dockerfile build.

SOLUTION

Add a benign comment after the command to invalidate the cache.

Starting with the Dockerfile in <https://github.com/docker-in-practice/todo> (which corresponds to the Step lines in the following output), we've done a build and then added a comment in the Dockerfile on the line with CMD. You can see the output of doing `docker build` again here:

```
$ docker build .
Sending build context to Docker daemon 2.56 kB
Sending build context to Docker daemon
Step 0 : FROM node
----> 91cbcf796c2c
Step 1 : MAINTAINER ian.miell@gmail.com
----> Using cache
----> 8f5a8a3d9240
Step 2 : RUN git clone -q https://github.com/docker-in-practice/todo.git
----> Using cache
----> 48db97331aa2
Step 3 : WORKDIR todo
----> Using cache
----> c5c85db751d6
Step 4 : RUN npm install
----> Using cache
----> be943c45c55b
Step 5 : EXPOSE 8000
----> Using cache
----> 805b18d28a65
Step 6 : CMD ["npm", "start"] #bust the cache
----> Running in fc6c4cd487ce
----> d66d9572115e
Removing intermediate container fc6c4cd487ce
Successfully built d66d9572115e
```

The reason this trick works is because Docker treats the non-whitespace change to the line as though it were a new command, so the cached layer is not re-used.

You may be wondering (as we did when we first looked at Docker) whether you can move Docker layers from image to image, merging them at will as though they were change sets in Git. This isn't possible at present within Docker. A layer is defined as a change set from a given image only. Because of this, once the cache has been broken, it can't be re-applied for commands re-used later in the build. For this reason, you're advised to put commands that are less likely to change nearer the top of the Dockerfile if possible.

DISCUSSION

For the initial iteration on a Dockerfile, splitting up every single command into a separate layer is excellent for speed of iteration, because you can selectively rerun parts of the process, as shown in the previous listing, but it's not so great for producing a small final image. It's not unheard-of for builds with a reasonable amount of complexity to approach the hard limit of 42 layers. To mitigate this, once you have a working build you're happy with, you should look at the steps in technique 56 for creating a production-ready image.

TECHNIQUE 23 **Intelligent cache-busting using build-args**

In the previous technique you saw how the cache can be busted mid-build by changing the relevant line.

In this technique we're going to take things a step further by controlling whether or not the cache is busted from the build command.

PROBLEM

You want to bust the cache on demand when performing a build, without editing the Dockerfile.

SOLUTION

Use the ARG directive in your Dockerfile to enable surgical cache-busting.

To demonstrate this, you're again going to use the Dockerfile at <https://github.com/docker-in-practice/todo>, but make a minor change to it.

What you want to do is control the busting of the cache before the `npm install`. Why would you want to do this? As you've learned, by default Docker will only break the cache if the command in the Dockerfile changes. But let's imagine that updated npm packages are available, and you want to make sure you get them. One option is to manually change the line (as you saw in the previous technique), but a more elegant way to achieve the same thing involves using the Docker ARG directive and a bash trick.

Add the ARG line to the Dockerfile as follows.

Listing 4.7 Simple Dockerfile with bustable cache

```
WORKDIR todo
ARG CACHEBUST=no
RUN npm install
```

← The ARG directive sets an environment variable for the build.

In this example, you use the ARG directive to set the CACHEBUST environment variable and default it to no if it's not set by the `docker build` command.

Now build that Dockerfile “normally”:

```
$ docker build .
Sending build context to Docker daemon 2.56kB
Step 1/7 : FROM node
latest: Pulling from library/node
aa18ad1a0d33: Pull complete
```



```

15a33158a136: Pull complete
f67323742a64: Pull complete
c4b45e832c38: Pull complete
f83e14495c19: Pull complete
41fea39113bf: Pull complete
f617216d7379: Pull complete
cbb91377826f: Pull complete
Digest: sha256:
➡ a8918e06476bef51ab83991aea7c199bb50bf131668c9739e6aa7984da1c1f6
Status: Downloaded newer image for node:latest
---> 9ealc3e33a0b
Step 2/7 : MAINTAINER ian.miell@gmail.com
---> Running in 03dba6770157
---> a5b55873d2d8
Removing intermediate container 03dba6770157
Step 3/7 : RUN git clone https://github.com/docker-in-practice/todo.git
---> Running in 23336fd5991f
Cloning into 'todo'...
---> 8ba06824d184
Removing intermediate container 23336fd5991f
Step 4/7 : WORKDIR todo
---> f322e2dbeb85
Removing intermediate container 2aa5ae19fa63
Step 5/7 : ARG CACHEBUST=no
---> Running in 9b4917f2e38b
---> f7e86497dd72
Removing intermediate container 9b4917f2e38b
Step 6/7 : RUN npm install
---> Running in a48e38987b04
npm info it worked if it ends with ok
[...]
added 249 packages in 49.418s
npm info ok
---> 324ba92563fd
Removing intermediate container a48e38987b04
Step 7/7 : CMD npm start
---> Running in ae76fa693697
---> b84dbc4bf5f1
Removing intermediate container ae76fa693697
Successfully built b84dbc4bf5f1

```

If you build it again with exactly the same `docker build` command, you'll observe that the Docker build cache is used, and no changes are made to the resulting image.

```

$ docker build .
Sending build context to Docker daemon 2.56kB
Step 1/7 : FROM node
---> 9ealc3e33a0b
Step 2/7 : MAINTAINER ian.miell@gmail.com
---> Using cache
---> a5b55873d2d8
Step 3/7 : RUN git clone https://github.com/docker-in-practice/todo.git
---> Using cache
---> 8ba06824d184

```

```

Step 4/7 : WORKDIR todo
----> Using cache
----> f322e2dbeb85
Step 5/7 : ARG CACHEBUST=no
----> Using cache
----> f7e86497dd72
Step 6/7 : RUN npm install
----> Using cache
----> 324ba92563fd
Step 7/7 : CMD npm start
----> Using cache
----> b84dbc4bf5f1
Successfully built b84dbc4bf5f1

```

At this point you decide that you want to force the npm packages to be rebuilt. Perhaps a bug has been fixed, or you want to be sure you're up to date. This is where the ARG variable you added to the Dockerfile in listing 4.7 comes in. If this ARG variable is set to a value never used before on your host, the cache will be busted from that point.

This is where you use the build-arg flag to docker build, along with a bash trick to force a fresh value:

```

$ docker build --build-arg CACHEBUST=${RANDOM} .
Sending build context to Docker daemon 4.096 kB
Step 1/9 : FROM node
----> 53d4d5f3b46e
Step 2/9 : MAINTAINER ian.miell@gmail.com
----> Using cache
----> 3a252318543d
Step 3/9 : RUN git clone https://github.com/docker-in-practice/todo.git
----> Using cache
----> c0f682653a4a
Step 4/9 : WORKDIR todo
----> Using cache
----> bd54f5d70700
Step 5/9 : ARG CACHEBUST=no
----> Using cache
----> 3229d52b7c33
Step 6/9 : RUN npm install
----> Running in 42f9b1f37a50
npm info it worked if it ends with ok
npm info using npm@4.1.2
npm info using node@v7.7.2
npm info attempt registry request try #1 at 11:25:55 AM
npm http request GET https://registry.npmjs.org/compression
npm info attempt registry request try #1 at 11:25:55 AM
[...]
Step 9/9 : CMD npm start
----> Running in 19219fe5307b
----> 129bab5e908a
Removing intermediate container 19219fe5307b
Successfully built 129bab5e908a

```

Run docker build with the build-arg flag, setting the CACHEBUST argument to a pseudo-random value generated by bash

Because the ARG CACHEBUST=no line itself has not changed, the cache is used here.

Because the CACHEBUST arg was set to a previously unset value, the cache is busted, and the npm install command is run again.

Note that the cache is busted on the line *following* the ARG line, not the ARG line itself. This can be a little confusing. The key thing to look out for is the “Running in” phrase—this means that a new container has been created to run the build line in.

The use of the `${RANDOM}` argument is worth explaining. Bash provides you with this reserved variable name to give you an easy way of getting a value between one and five digits long:

```
$ echo ${RANDOM}
19856
$ echo ${RANDOM}
26429
$ echo ${RANDOM}
2856
```

This can come in handy, such as when you want a probably unique value to create files just for a specific run of a script.

You can even produce a much longer random number if you’re concerned about clashes:

```
$ echo ${RANDOM}${RANDOM}
434320509
$ echo ${RANDOM}${RANDOM}
1327340
```

Note that if you’re not using bash (or a shell that has this `RANDOM` variable available), this technique won’t work. In that case, you could use the `date` command instead to produce a fresh value:

```
$ docker build --build-arg CACHEBUST=$(date +%s) .
```

DISCUSSION

This technique has demonstrated a few things that will come in handy when using Docker. You’ve learned about using the `--build-args` flag to pass in a value to the Dockerfile and bust the cache on demand, creating a fresh build without changing the Dockerfile.

If you use bash, you’ve also learned about the `RANDOM` variable, and how it can be useful in other contexts than just Docker builds.

TECHNIQUE 24 **Intelligent cache-busting using the ADD directive**

In the previous technique you saw how the cache could be busted mid-build at a time of your choosing, which was itself a level up from using the `--no-cache` flag to ignore the cache completely.

Now you’re going to take it to the next level, so that you can automatically bust the cache only when it’s necessary to. This can save you a lot of time and compute—and therefore money!

PROBLEM

You want to bust the cache when a remote resource has changed.

SOLUTION

Use the Dockerfile `ADD` directive to only bust the cache when the response from a URL changes.

One of the early criticisms of Dockerfiles was that their claim of producing reliable build results was misleading. Indeed, we took this very subject up with the creator of Docker back in 2013 (<http://mng.bz/B8E4>).

Specifically, if you make a call to the network with a directive in your Dockerfile like this,

```
RUN git clone https://github.com/nodejs/node
```

then by default the Docker build will perform this once per Docker daemon. The code on GitHub could change substantially, but as far as your Docker daemon is concerned, the build is up to date. Years could pass, and the same Docker daemon will still be using the cache.

This may sound like a theoretical concern, but it's a very real one for many users. We've seen this happen many times at work, causing confusion. You've already seen some solutions to this, but for many complex or large builds, those solutions are not granular enough.

THE SMART CACHE-BUSTING PATTERN

Imagine you have a Dockerfile that looks like the following listing (note that it won't work! It's just a Dockerfile pattern to show the principle).

Listing 4.8 An example Dockerfile

Installs a series of packages as a prerequisite

```
FROM ubuntu:16.04
→ RUN apt-get install -y git and many other packages
  RUN git clone https://github.com/nodejs/node
  WORKDIR node
  RUN make && make install
```

Clones a regularly changing repository (nodejs is just an example)

Runs a make and install command, which builds the project

This Dockerfile presents some challenges to creating an efficient build process. If you want to build everything from scratch each time, the solution is simple: use the `--no-cache` argument to `docker build`. The problem with this is that each time you run a build you're repeating the package installation in the second line, which is (mostly) unnecessary.

This challenge can be solved by busting the cache just before the `git clone` (demonstrated in the last technique). This raises another challenge, however: what if the Git repository hasn't changed? Then you're doing a potentially costly network transfer, followed by a potentially costly `make` command. Network, compute, and disk resources are all being used unnecessarily.

One way to get around this is to use technique 23, where you pass in a build argument with a new value every time you know that the remote repository has changed. But this still requires manual investigation to determine whether there has been a change, and intervention.

What you need is a command that can determine whether the resource has changed since the last build, and only then bust the cache.

THE ADD DIRECTIVE—UNEXPECTED BENEFITS

Enter the ADD directive!

You're already familiar with ADD, as it's a basic Dockerfile directive. Normally it's used to add a file to the resulting image, but there are two useful features of ADD that you can use to your advantage in this context: it caches the contents of the file it refers to, and it can take a network resource as an argument. This means that you can bust the cache whenever the output of a web request changes.

How can you take advantage of this when cloning a repository? Well, that depends on the nature of the resource you're referencing over the network. Many resources will have a page that changes when the repository itself changes, but these will vary from resource type to resource type. Here we'll focus on GitHub repos, because that's a common use case.

The GitHub API provides a useful resource that can help here. It has URLs for each repository that return JSON for the most recent commits. When a new commit is made, the content of the response changes.

Listing 4.9 Using ADD to trigger a cache bust

It doesn't matter where the output of the file goes, so we send it to /dev/null.

```
FROM ubuntu:16.04
ADD https://api.github.com/repos/nodejs/node/commits
➡ /dev/null
RUN git clone https://github.com/nodejs/node
[...]
```

The URL that changes when a new commit is made

The git clone will take place only when a change is made.

The result of the preceding listing is that the cache is busted only when a commit has been made to the repo since the last build. No human intervention is required, and no manual checking.

If you want to test this mechanism with a frequently changing repo, try using the Linux kernel.

Listing 4.10 Adding the Linux kernel code to an image

```
FROM ubuntu:16.04
ADD https://api.github.com/repos/torvalds/linux/commits /dev/null
RUN echo "Built at: $(date)" >> /build_time
```

The ADD command, this time using the Linux repository

Outputs the system date into the built image, which will show when the last cache-busting build took place

If you create a folder and put the preceding code into a Dockerfile, and then run the following command regularly (every hour, for example), the output date will change only when the Linux Git repo changes.

Listing 4.11 Building a Linux code image

```
$ docker build -t linux_last_updated .
$ docker run linux_last_updated cat /build_time
```

Builds the image and gives it the name `linux_last_updated`

Outputs the contents of the `/build_time` file from the resulting image

DISCUSSION

This technique demonstrated a valuable automated technique for ensuring builds only take place when necessary.

It also demonstrated some of the details of how the `ADD` command works. You saw that the “file” could be a network resource, and that if the contents of the file (or network resource) change from a previous build, a cache bust takes place.

In addition, you also saw that network resources have related resources that can indicate whether the resource you’re referencing has changed. Although you could, for example, reference the main GitHub page to see if there are any changes there, it’s likely that the page changes more frequently than the last commit (such as if the time of the web response is buried in the page source, or if there’s a unique reference string in each response).

In the case of GitHub, you can reference the API, as you saw. Other services, such as BitBucket, offer similar resources. The Kubernetes project, for example, offers this URL to indicate which release is stable: <https://storage.googleapis.com/kubernetes-release/release/stable.txt>. If you were building a Kubernetes-based project, you might put an `ADD` line in your Dockerfile to bust the cache whenever this response changes.

TECHNIQUE 25 Setting the right time zone in your containers

If you’ve ever installed a full operating system, you’ll know that setting the time zone is part of the process of setting it up. Even though a container isn’t an operating system (or a virtual machine), it contains the files that tell programs how to interpret the time for the configured timezone.

PROBLEM

You want to set the time zone correctly for your containers.

SOLUTION

Replace the container’s localtime file with a link to the time zone you want.

The following listing demonstrates the problem. It doesn’t matter where in the world you run this, the container will show the same time zone.

Listing 4.12 Container starting with wrong time zone

```

The time zone on the host is GMT.  $ date +%Z
                                  GMT
                                  $ docker run centos:7 date +%Z
The time zone in the container is GMT.  UTC

```

Runs a command to display the time zone on the host

Runs a container and outputs the date from within that

The container contains the files that determine which time zone is used by the container to interpret the time value it gets. The actual time used is, of course, tracked by the host operating system.

The next listing shows how you can set the time zone to the one you want.

Listing 4.13 Dockerfile for replacing the centos:7 default time zone

```

Removes the existing localtime symlink file  FROM centos:7
                                                Starts from the centos image we just looked at
RUN rm -rf /etc/localtime
RUN ln -s /usr/share/zoneinfo/GMT /etc/localtime  Replaces the /etc/localtime link with a link to the time zone you want
CMD date +%Z
Shows the time zone of your container as the default command to run

```

In listing 4.13 the key file is `/etc/localtime`. This points to the file that tells the container which time zone to use when it's asked for the time. The default time given is in the UTC time standard, which is used if the file doesn't exist (the minimal BusyBox image, for example, doesn't have it).

The following listing shows the output of building the preceding Dockerfile.

Listing 4.14 Building a time-zone-replacing Dockerfile

```

$ docker build -t timezone_change .
Sending build context to Docker daemon 62.98 kB
Step 1 : FROM centos:7
7: Pulling from library/centos
45a2e645736c: Pull complete
Digest: sha256:
c577af3197aacedf79c5a204cd7f493c8e07ffbce7f88f7600bf19c688c38799
Status: Downloaded newer image for centos:7
---> 67591570dd29
Step 2 : RUN rm -rf /etc/localtime
---> Running in fb52293849db
---> 0deda41be8e3
Removing intermediate container fb52293849db
Step 3 : RUN ln -s /usr/share/zoneinfo/GMT /etc/localtime
---> Running in 47bf21053b53

```

Builds the container

```

----> 5b5cb1197183
Removing intermediate container 47bf21053b53
Step 4 : CMD date +%Z
----> Running in 1e481eda8579
----> 9477cdaa73ac
Removing intermediate container 1e481eda8579
Successfully built 9477cdaa73ac
$ docker run timezone_change          ← Runs the container
GMT                                  ← Outputs the specified time zone

```

In this way you can specify the time zone to use within—and only within—your container. Many applications depend on this setting, so it comes up not infrequently if you’re running a Docker service.

There’s another problem that this container-level time granularity can solve. If you’re working for a multinational organization and run many different applications on servers based in data centers around the world, the ability to change the time zone in your image and trust that it will report the right time wherever it lands is a useful trick to have to hand.

DISCUSSION

Because the point of Docker images is explicitly to provide a consistent experience no matter where you run your container, there are a number of things you may stumble across if you do want varied results depending on where an image is deployed.

As an example, if you’re automatically producing CSV spreadsheets of data for users in different locations, they may have certain expectations of the data format. American users might expect dates in the mm/dd format, whereas Europeans might expect dates in dd/mm format, and Chinese users might expect dates in their own character set.

In the next technique we’ll consider locale settings, which affect how dates and times are printed in the `local` format, among other things.

TECHNIQUE 26 **Locale management**

In addition to time zones, locales are another aspect of Docker images that can be relevant when building images or running containers.

NOTE A locale defines which language and country settings your programs should use. Typically a locale will be set in the environment through the `LANG`, `LANGUAGE`, and `locale-gen` variables, and through variables beginning with `LC_`, such as `LC_TIME`, whose setting determines how the time is displayed to the user.

NOTE An encoding (in this context) is the means by which text is stored as bytes on a computer. A good introduction to this subject is available from W3C here: <https://www.w3.org/International/questions/qa-what-is-encoding>. It’s worth taking the time to understand this subject, as it comes up in all sorts of contexts.

PROBLEM

You're seeing encoding errors in your application builds or deployments.

SOLUTION

Ensure the language-specific environment variables are correctly set in your Dockerfile.

Encoding issues aren't always obvious to all users, but they can be fatal when building applications.

Here are a couple of examples of typical encoding errors when building applications in Docker.

Listing 4.15 Typical encoding errors

```
MyFileDialog:66: error: unmapable character for encoding ASCII
```

```
UnicodeEncodeError: 'ascii' codec can't encode character u'\xa0' in  
↳ position 20: ordinal not in range(128)
```

These errors can kill a build or an application stone dead.

TIP A non-exhaustive list of key words to look out for in the error are “encoding,” “ascii,” “unicode,” “UTF-8,” “character,” and “codec.” If you see these words, chances are you're dealing with an encoding issue.

WHAT DOES THIS HAVE TO DO WITH DOCKER?

When you set up a full-blown operating system, you're typically guided through a setup process that asks you to confirm your preferred time zone, language, keyboard layout, and so on.

Docker containers, as you know by now, aren't full-blown operating systems set up for general use. Rather they're (increasingly) minimal environments for running applications. By default, therefore, they may not come with all the setup you're used to with an operating system.

In particular, Debian removed their dependency on the locales package in 2011, which means that, by default, there's no locale setup in a container based on a Debian image. For example, the following listing shows a Debian-derived Ubuntu image's default environment.

Listing 4.16 Default environment on an Ubuntu container

```
$ docker run -ti ubuntu bash  
root@d17673300830:/# env  
HOSTNAME=d17673300830  
TERM=xterm  
LS_COLORS=rs=0 [...]  
HIST_FILE=/root/.bash_history  
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin  
PWD=/  
SHLV=1  
HOME=/root  
_=/usr/bin/envj
```

There are no `LANG` or similar `LC_` settings available in the image by default.

Our Docker host is shown in the next listing.

Listing 4.17 LANG setting on Docker host OS

```
$ env | grep LANG
LANG=en_GB.UTF-8
```

There's a `LANG` setting in our shell that informs applications that the preferred encoding in our terminal is British English, with text encoded in UTF-8.

To demonstrate an encoding issue, we'll create a file locally that contains a UTF-8-encoded UK currency symbol (the UK's pound sign), and then show how the interpretation of that file changes depending on the terminal's encoding.

Listing 4.18 Creating and showing a UTF-8-encoded UK currency symbol

```
$ env | grep LANG
LANG=en_GB.UTF-8
$ echo -e "\xc2\xa3" > /tmp/encoding_demo
$ cat /tmp/encoding_demo
£
```

Uses `echo` with the `-e` flag to output two bytes into a file, which represent a UK pound sign

Cats the file; we'll see a pound sign.

In UTF-8, a pound sign is represented by two bytes. We output these two bytes using `echo -e` and the `\x` notation and redirect output into a file. When we `cat` the file, the terminal reads the two bytes and knows to interpret the output as a pound sign.

Now if we change our terminal's encoding to use the Western (ISO Latin 1) encoding (which sets up our local `LANG` also) and output the file, it looks quite different:

Listing 4.19 Demonstrating the encoding problem with the UK currency symbol

```
$ env | grep LANG
LANG=en_GB.ISO8859-1
$ cat /tmp/encoding_demo
Â£
```

The `LANG` environment variable is now set to Western (ISO Latin 1), which is set by the terminal.

The two bytes are interpreted differently, as two separate characters that are displayed to us.

The `\xc2` byte is interpreted as a capital *A* with a circumflex on top, and the `\xa3` byte is interpreted as a UK pound sign!

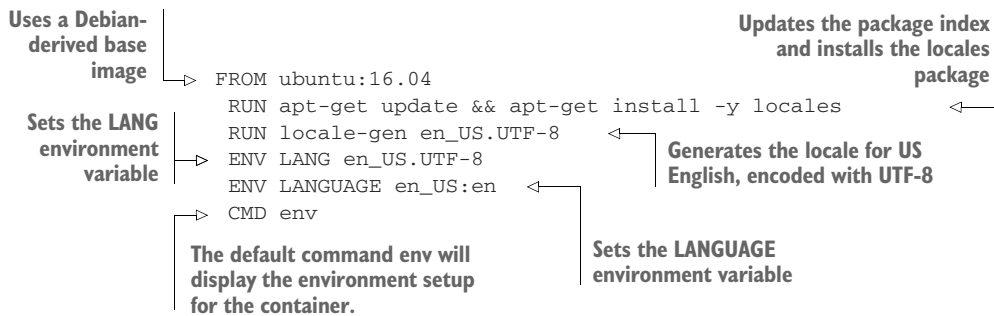
NOTE We say “we” rather than “you” above deliberately! Debugging and controlling encodings is a tricky affair, which can depend on a combination of the running application's state, the environment variables you have set up, the running application, and all of the preceding factors that create the data you're examining!

As you’ve seen, encodings can be affected by the encoding set in the terminal. Getting back to Docker, we noted that no encoding environment variables were set by default in our Ubuntu container. Because of this, you can get different results when running the same commands on your host or in a container. If you see errors that seem to relate to encodings, you may need to set them in your Dockerfile.

SETTING UP ENCODINGS IN A DOCKERFILE

We’ll now look at how you can control the encoding of a Debian-based image. We’ve chosen this image because it’s likely to be one of the more common contexts. This example will set up a simple image that just outputs its default environment variables.

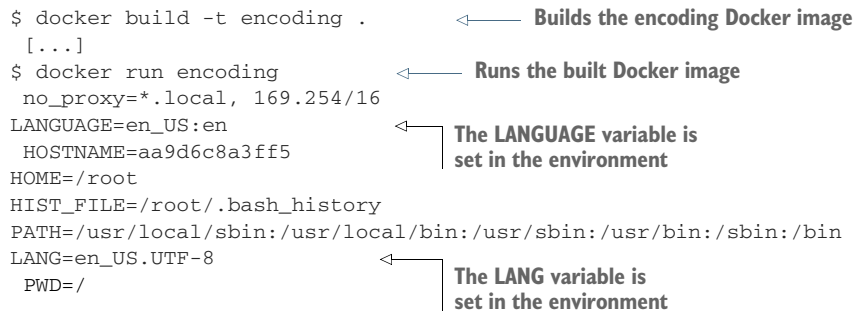
Listing 4.20 Setting up a Dockerfile example



You may be wondering what the differences between the `LANG` and `LANGUAGE` variables are. Briefly, `LANG` is the default setting for the preferred language and encoding settings. It also provides a default when applications look for the more specific `LC_*` settings. `LANGUAGE` is used to provide an ordered list of languages preferred by applications if the principal one isn’t available. More information can be found by running `man locale`.

Now you can build the image, and run it to see what’s changed.

Listing 4.21 Building and running the encoding image



DISCUSSION

Like the previous time zone technique, this technique illustrates an issue that catches people out on a regular basis. Like many of the more irritating issues we come across, these don't always make themselves obvious when the image is being built, which makes the time wasted debugging these issues very frustrating. For this reason, it's worth keeping these settings in mind when supporting others who are using Docker images.

TECHNIQUE 27 **Stepping through layers with the image-stepper**

If you've built an image that has a number of steps, you can often find yourself in the position of wanting to know where a particular file was introduced, or what state it was in at a particular point in the build. Combing through each image layer can be laborious, because you'll have to determine the order of layers, retrieve each ID, and start each one up using that ID.

This technique shows you a one-liner that tags each layer of the build in order, meaning you only have to increment a number to work through the images and find out whatever it is you need to know.

PROBLEM

You want to easily refer to each step of your build.

SOLUTION

Use the `docker-in-practice/image-stepper` image to order the tags for your image.

To illustrate this technique, we'll first show you a script that achieves this result so you understand how it works. Then we'll give you a constructed image to make achieving the result easier.

Here's a simple script that tags every layer in a given image (`myimage`) in the order of creation.

The Dockerfile for `myimage` follows.

Listing 4.22 Dockerfile for image with multiple layers

```
FROM debian
RUN touch /file1
RUN touch /file2
RUN touch /file3
RUN touch /file4
RUN touch /file5
RUN touch /file6
RUN touch /file7
RUN touch /file8
RUN touch /file9
RUN touch /file10
CMD ["cat", "/file1"]
```

← Uses debian as a base image

Creates 10 files in separate layers

← Runs a bespoke command that cats the first file

This is a simple enough Dockerfile, but it's one where it will be clear which stage you are at in the build.

Build this docker image with the following command.

Listing 4.23 Building the myimage image

```
$ docker build -t myimage -q .
sha256:b21d1e1da994952d8e309281d6a3e3d14c376f9a02b0dd2ecbe6cabffea95288
```

Builds the image with the quiet (-q) flag, tagging it as myimage

The image identifier is the only output.

Once the image is built, you can run the following script.

Listing 4.24 Tagging each layer of myimage in numerical order

```
#!/bin/bash
x=1
for id in $(docker history -q "myimage:latest" |
  grep -vw missing
  | tac)
do
  docker tag "${id}" "myimage:latest_step_${x}"
  ((x++))
done
```

Doesn't consider the remotely built layers, which are marked as missing (see the note below)

Initializes the counter variable (x) to 1

Runs a for loop to retrieve the history of the image

Uses the tac utility to reverse the order of image IDs the docker history command outputs

In each iteration of the loop, tags the image appropriately with the incrementing number

Increments the step counter

If you save the preceding file as tag.sh and run it, the image will be tagged in layer order.

NOTE This tagging method technique will only work on images built locally. See the note in technique 16 for more information.

Listing 4.25 Tagging and showing the layers

```
$ ./tag.sh
$ docker images | grep latest_step
myimage latest_step_12 1bfca0ef799d 3 minutes ago 123.1 MB
myimage latest_step_11 4d7f66939a4c 3 minutes ago 123.1 MB
myimage latest_step_10 78d31766b5cb 3 minutes ago 123.1 MB
myimage latest_step_9 f7b4dcbbd74f 3 minutes ago 123.1 MB
myimage latest_step_8 69b2fa0ce520 3 minutes ago 123.1 MB
myimage latest_step_7 b949d71fb58a 3 minutes ago 123.1 MB
myimage latest_step_6 8af3bbf1e7a8 3 minutes ago 123.1 MB
myimage latest_step_5 ce3dfbdfed74 3 minutes ago 123.1 MB
myimage latest_step_4 598ed62cabb9 3 minutes ago 123.1 MB
myimage latest_step_3 6b290f68d4d5 3 minutes ago 123.1 MB
myimage latest_step_2 586da987f40f 3 minutes ago 123.1 MB
myimage latest_step_1 19134a8202e7 7 days ago 123.1 MB
```

Runs the script from listing 4.24

Runs a docker images command with a simple grep to see the tagged layers

The original (and older) base image has also been tagged as latest_step_1.

The steps to build the myimage image

Now that you've seen the principle, we'll demonstrate how to dockerize this one-off script and make it work for the general case.

NOTE The code for this technique is available at <https://github.com/docker-in-practice/image-stepper>.

First, turn the previous script into a script that can take arguments.

Listing 4.26 Generic tagging script for the image-stepper image

```
#!/bin/bash
IMAGE_NAME=$1
IMAGE_TAG=$2
if [[ $IMAGE_NAME = '' ]]
then
    echo "Usage: $0 IMAGE_NAME [ TAG ]"
    exit 1
fi
if [[ $IMAGE_TAG = '' ]]
then
    IMAGE_TAG=latest
fi
x=1
for id in $(docker history -q "${IMAGE_NAME}:${IMAGE_TAG}" |
➡ grep -vw missing | tac)
do
    docker tag "${id}" "${IMAGE_NAME}:${IMAGE_TAG}_step_${x}"
    ((x++))
done
```

Defines a bash script that can take two arguments: the image name to process, and the tag you want to step up to

The script from listing 4.24, with the arguments substituted in

You can then embed the script in listing 4.26 into a Docker image that you place into a Dockerfile and run as the default ENTRYPOINT.

Listing 4.27 Dockerfile for image-stepper image

```
FROM ubuntu:16.04
RUN apt-get update -y && apt-get install -y docker.io
ADD image_stepper /usr/local/bin/image_stepper
ENTRYPOINT ["/usr/local/bin/image_stepper"]
```

Adds the script from listing 4.26 to the image

Uses Ubuntu as a base layer

Installs docker.io to get the Docker client binary

Runs the image_stepper script by default

The Dockerfile in listing 4.27 creates an image that runs the script in listing 4.26. The command in listing 4.28 runs this image, giving `myimage` as an argument.

This image, when run against another Docker image built on your host, will then create the tags for each step, allowing you to easily look at the layers in order.

The version of the client binary installed by the `docker.io` package must be compatible with the version of the Docker daemon on your host machine, typically meaning the client must not be newer.

Listing 4.28 Running image-stepper against another image

Runs the image-stepper image as a container, and removes the container when done

Mounts the host's docker socket, so you can use the Docker client installed in listing 4.27

Downloads the image-stepper image from the Docker Hub

```
$ docker run --rm
➤ -v /var/run/docker.sock:/var/run/docker.sock
➤ dockerinpractice/image-stepper
➤ myimage
```

Tags the myimage created previously

```
Unable to find image 'dockerinpractice/image-stepper:latest' locally
latest: Pulling from dockerinpractice/image-stepper
b3e1c725a85f: Pull complete
4daad8bdde31: Pull complete
63fe8c0068a8: Pull complete
4a70713c436f: Pull complete
bd842a2105a8: Pull complete
1a3a96204b4b: Pull complete
d3959cd7b55e: Pull complete
Digest: sha256:
➤ 65e22f8a82f2221c846c92f72923927402766b3c1f7d0ca851ad418fb998a753
Status: Downloaded newer image for dockerinpractice/image-stepper:latest
$ docker images | grep myimage
```

The output of the docker run command

myimage	latest	2c182dabe85c	24 minutes ago	123 MB
myimage	latest_step_12	2c182dabe85c	24 minutes ago	123 MB
myimage	latest_step_11	e0ff97533768	24 minutes ago	123 MB
myimage	latest_step_10	f46947065166	24 minutes ago	123 MB
myimage	latest_step_9	8a9805a19984	24 minutes ago	123 MB
myimage	latest_step_8	88e42bed92ce	24 minutes ago	123 MB
myimage	latest_step_7	5e638f955e4a	24 minutes ago	123 MB
myimage	latest_step_6	f66b1d9e9cbd	24 minutes ago	123 MB
myimage	latest_step_5	bd07d425bd0d	24 minutes ago	123 MB
myimage	latest_step_4	ba913e75a0b1	24 minutes ago	123 MB
myimage	latest_step_3	2ebcda8cd503	24 minutes ago	123 MB
myimage	latest_step_2	58f4ed4fe9dd	24 minutes ago	123 MB
myimage	latest_step_1	19134a8202e7	2 weeks ago	123 MB

Runs docker images and greps out the images you've just tagged

```
$ docker run myimage:latest_step_8 ls / | grep file
```

The files shown are those created up to that step.

```
file1
file2
file3
file4
file5
file6
file7
```

The images are tagged.

Picks a step at random and lists the files in the root directory, grep-ing out the ones created in the Dockerfile from listing 4.27

NOTE On non-Linux OSs (such as Mac and Windows) you may need to specify the folder in which Docker runs in your Docker preferences as a file sharing setting.

This technique is useful for seeing where a particular file was added within a build, or what state a file was in at a particular point in the build. When debugging a build, this can be invaluable!

DISCUSSION

This technique is used in technique 52 to demonstrate that a deleted secret is accessible within a layer within an image.

TECHNIQUE 28 **Onbuild and golang**

The `ONBUILD` directive can cause a lot of confusion for new Docker users. This technique demonstrates its use in a real-world context by building and running a Go application with a two-line Dockerfile.

PROBLEM

You want to reduce the steps in building an image required for an application.

SOLUTION

Use the `ONBUILD` command to automate and encapsulate the building of an image.

First you'll run the process through, and then we'll explain what's going on. The example we'll use is the `outyet` project, which is an example in the `golang` GitHub repository. All it does is set up a web service that returns a page telling you whether Go 1.4 is available yet.

Build the image as follows.

Listing 4.29 Building the outyet image

Navigates to the outyet folder	<pre>\$ git clone https://github.com/golang/example</pre>	← Clones the Git repository
	<pre>\$ cd example/outyet</pre>	
	<pre>\$ docker build -t outyet .</pre>	← Builds the outyet image

Run a container from the resulting image, and retrieve the served web page.

Listing 4.30 Running and validating the outyet image

<p>The <code>--publish</code> flag tells Docker to publish the container's port 8080 on the external port 6060.</p>	<p>The <code>--name</code> flag gives your container a predictable name to make it easier to work with.</p>	<p>Runs the container in the background</p>	<p>← Curls the output container's port</p>
<pre>\$ docker run</pre>			
<pre> ➡ --publish 8080:8080</pre>			
<pre> ➡ --name outyet1 -d outyet</pre>			
<pre>\$ curl localhost:8080</pre>			
<pre><!DOCTYPE html><html><body><center></pre>			
<pre> <h2>Is Go 1.4 out yet?</h2></pre>			
<pre> <h1></pre>			
<pre> YES!</pre>			
<pre> </h1></pre>			
<pre></center></body></html></pre>			

**The
webpage
that the
container
serves**

That's it—a simple application that returns a web page that tells you whether Go 1.4 is out yet or not.

If you look around the cloned repository, you'll see the Dockerfile is just two lines!

Listing 4.31 The onyet Dockerfile

```
FROM golang:onbuild
EXPOSE 8080
```

Starts the build from the golang:onbuild image

Exposes port 8080

Confused yet? OK, it may make more sense when you look at the Dockerfile for the golang:onbuild image.

Listing 4.32 The golang:onbuild Dockerfile

```
FROM golang:1.7
RUN mkdir -p /go/src/app
WORKDIR /go/src/app
CMD ["go-wrapper", "run"]
ONBUILD COPY . /go/src/app
ONBUILD RUN go-wrapper download
ONBUILD RUN go-wrapper install
```

Sets the resulting image's command to call the go-wrapper to run the go app

Uses the golang:1.7 image as a base

Makes a folder to store the application in

Moves into that folder

The first ONBUILD command copies the code in the context of the Dockerfile into the image.

The second ONBUILD command downloads any dependencies, again using the go-wrapper command.

The third ONBUILD

The golang:onbuild image defines what happens when the image is used in the FROM directive in any other Dockerfile. The result is that when a Dockerfile uses this image as a base, the ONBUILD commands will fire as soon as the FROM image is downloaded, and (if not overridden) the CMD will be run when the resulting image is run as a container.

Now the output of the `docker build` command in the next listing may make more sense.

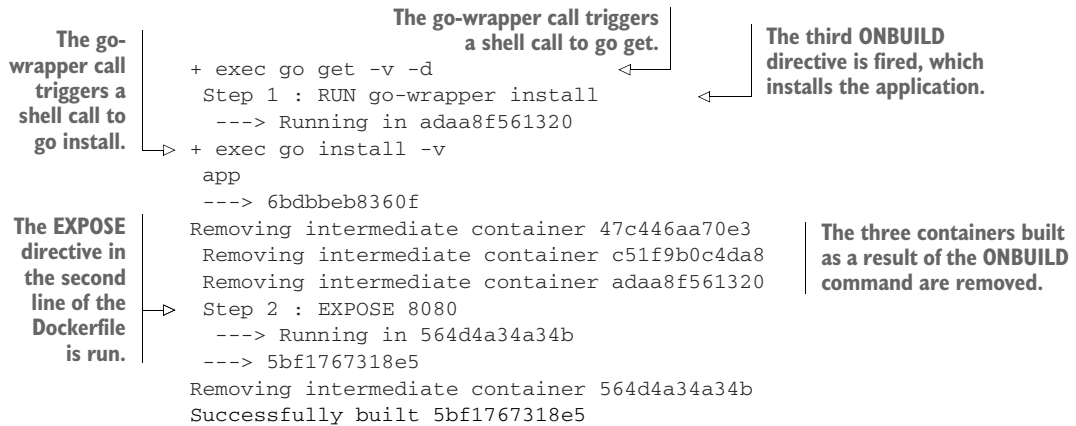
```
Step 1 : FROM golang:onbuild
onbuild: Pulling from library/golang
6d827a3ef358: Pull complete
2726297beaf1: Pull complete
7d27bd3d7fec: Pull complete
62ace0d726fe: Pull complete
af8d7704cf0d: Pull complete
6d8851391f39: Pull complete
988b98d9451c: Pull complete
5bbc96f59ddc: Pull complete
Digest: sha256:
886a63b8de95d5767e779dee4ce5ce3c0437fa48524aedd93199fb12526f15e0
Status: Downloaded newer image for golang:onbuild
# Executing 3 build triggers...
Step 1 : COPY . /go/src/app
Step 1 : RUN go-wrapper download
--> Running in c51f9b0c4da8
```

The Docker build signals its intention to run the ONBUILD directives.

The second ONBUILD directive is fired, which downloads.

The FROM directive is run, and the golang:onbuild image is

The first ONBUILD directive copies the Go code in the Dockerfile's context into the build.



The result of this technique is that you have an easy way to build an image that only contains the code required to run it, and no more. Leaving the build tools lying around in the image not only makes it larger than it needs to be, but also increases the security attack surface of the running container.

DISCUSSION

Because Docker and Go are fashionable technologies currently often seen together, we've used this to demonstrate how `ONBUILD` can be used to build a Go binary.

Other examples of `ONBUILD` images exist. There are `node:onbuild` and `python:onbuild` images available on Docker Hub.

It's hoped that this might inspire you to construct your own `ONBUILD` image that could help your organization with common patterns of building. This standardization can help reduce impedance mismatch between different teams even further.

Summary

- You can insert files from your local machine and from the internet into images.
- The cache is a crucial part of building images, but it can be a fickle friend and occasionally needs prompting to do what you want.
- You can “bust” the cache using build arguments or using the `ADD` directive, or you can ignore the cache completely with the `no-cache` option.
- The `ADD` directive is generally used to inject local files and folders into the built image.
- System configuration may still be relevant inside Docker, and image build time is a great time to do it.
- You can debug your build process using the “image-stepper” technique (technique 27), which tags each stage of the build for you.
- The time zone setting is the most common “gotcha” when configuring containers, especially when you are a non-U.S. or multinational company.
- Images with `ONBUILD` are very easy to use, because you might not need to customize the build at all.